

How We Run AI-Assisted Engineering


The operating model


The map the other talks slot into.

The trap

- "Use AI to code" sounds like **one tool** — a model, a clever prompt.
- Run it at **team scale** and that breaks: collisions, forgotten context, unreviewed output, no consistency.

“ What actually works is an **operating model** — a small set of layers around how you already ship. ”





The loop in motion

Intake → spec → plan → dispatch → **isolated build** → **adversarial review** → merge → **write memory** → recall feeds the next task.

The four layers – each its own talk

- **Authored rules** – the versioned constitution the agent reads every time. → *The Repo Is the Onboarding Doc*
- **Learned memory** – curated, tiered facts that compound. → *Tiered Agent Memory*
- **Live coordination** – many agents in parallel, no collision. → *The Virtual Engineering Team*
- **Adversarial verification** – cross-vendor review + evidence. → *Trust, but Verify*

Wrapped in governance

Autonomy is **risk-tiered**, not all-or-nothing:

- cheap, reversible work runs **autonomously** to a PR
- expensive or irreversible work hits a **human gate**

→ companion: *How Much Leash?*

What it buys

The SDLC stays familiar to every engineer — the **layers** do the work:

- **Trustworthy** (verification)
- **Parallel** (coordination)
- **Consistent** (rules)
- **Compounding** (memory)

“ *Honest caveat:* the operating layer is real overhead. It earns its keep **at scale**, not on a one-off script. ”

Close

"Not one tool – a thin operating system around how you already ship software."

Five short talks, one system. This was the map.