

How Much Leash?

Risk-tiered autonomy

Match the leash to the blast radius.

The false binary

- The usual debate: **"let agents run free"** vs **"approve everything."**
- One is **reckless**. The other **throws away the speed** that made agents worth using.

“ The way out: stop treating autonomy as a global switch. Make it a **per-task routing decision**. ”


The axis that matters


Not "*how hard is the task?*" — but:

"How bad if it's wrong, and can we undo it?"

Blast radius × reversibility.

- Hard task, tiny blast radius → safe to automate.
- Trivial task, irreversible → **not**.





The gate is lightweight

- A human confirms the **plan and direction** – not every keystroke.
- Approve the **approach**, then let the agent execute.

“ Heavy gates get rubber-stamped or bypassed. A cheap gate is a gate that actually gets used.”

The autonomous lane is still safe

- "Autonomous" means **to a PR — not to prod.**
- Review and **revert** still apply. A human is **notified.**

|“ Autonomy ≠ unsupervised merge to main.

”

Re-tier dynamically

- A task that **surprises** — touches more than expected, fails a check — **escalates** to a gate.
- **Downgrade trust on anomaly;** don't ride the original tier blindly.

This is your world already

Risk-tiered autonomy is just **change-management / blast-radius policy** — applied to agents.

The same instinct behind *"who's allowed to approve a prod migration?"*

What it buys

Speed where it's safe. Control where it counts.

“ *Honest caveat:* the risk is **mis-tiering**. Err toward the gate when unsure, and keep the reversibility assessment honest — the cost of undo is easy to underestimate. ”

Close

"Don't ask 'should agents be autonomous?' Ask 'how much leash for *this* task?'"