

Tiered Agent Memory

Giving AI coding agents an institutional memory

A context-engineering pattern we landed on.

The two problems


- AI agents are **amnesiac** — every session starts from zero, so the same constraints, preferences, and past mistakes get re-explained endlessly.
- Context is **finite and metered** — you *can't* paste everything you know into every session: it's expensive, slow, and dilutes the model's attention on the actual task.

“ The tension: more knowledge available, but loading more makes the agent **worse and costlier**. ”

The insight

The same problem **hardware solved decades ago**: you can't keep everything fast and close — so you **tier** it.

small / hot / always-on → larger / warm / on-demand →
large / cold / fetched-when-relevant



Pillar 1 — Atomic, typed facts

- **One fact = one file.** Atomicity makes facts easy to update, delete, and dedupe — no giant doc nobody maintains.
- Each fact is **typed**: who-the-user-is • feedback/corrections (*with the why*) • ongoing-work context • external references.
- Type drives **policy**: how long a fact lives, how aggressively it's recalled.

Pillar 2 – Tiered loading *(the cost optimization)*

- **Hot index** – one line per fact, always loaded. Just enough to know a fact exists and when it's relevant. Deliberately tiny, cross-cutting only.
- **Warm catalogs** – domain-grouped, loaded only when the task enters that domain.
- **Cold recall** – full detail pulled in for a single fact only when its summary matches.

“ This is where the **token + attention savings** come from. ”


Pillar 3 — Relevance-based recall

- Each fact carries a one-line **description written for matching**, not for humans.
- The system surfaces **only relevant facts**...
- ...and surfaces them as **background context, not commands** — so old notes can't hijack new work.

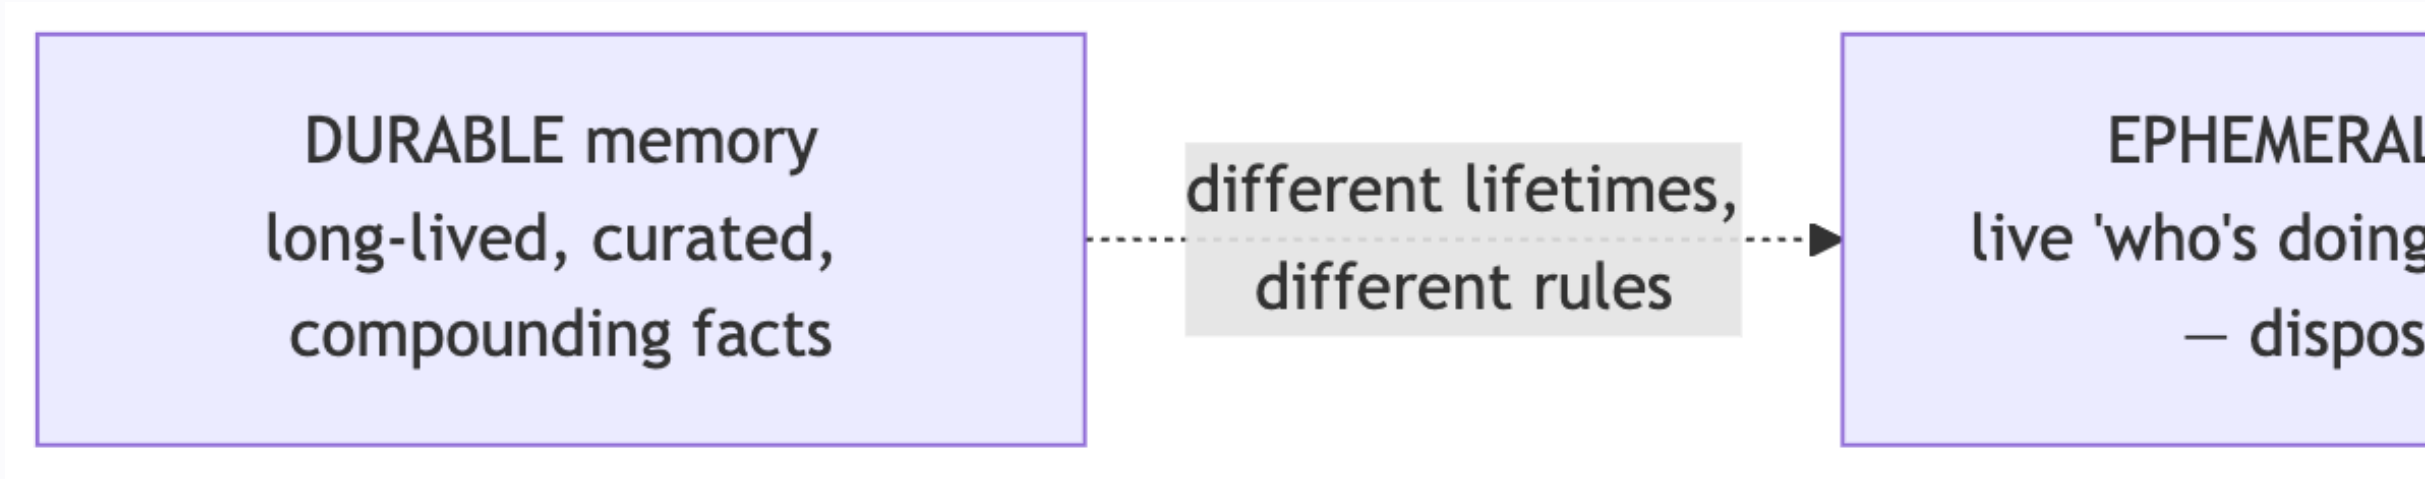
Pillar 4 — Curation discipline (*why it doesn't rot*)

The failure mode of every knowledge base is **staleness**. Four rules prevent it:

- **Dedup** — update the existing fact, never add a near-duplicate.
- **Delete-on-wrong** — a disproven fact is removed, not buried.
- **Single source of truth** — don't store what code/git/docs already record; point to it.
- **Verify-on-use + absolute dates** — a memory reflects what was true *when written*; confirm before acting, never store "yesterday."



Pillar 5 – Separate durable vs ephemeral



- Long-lived knowledge and live "who's-doing-what-now" coordination have **different lifetimes and different rules**.
- Mixing them is how knowledge bases fill with **noise**. Keep two channels.

What changes in practice

- **Knowledge compounds** — a correction is captured once, applied every future session.
- **Lower context cost** — always-on footprint stays tiny vs. dumping everything.
- **More consistent behavior** — fewer repeated mistakes, less re-explaining.
- **Trustworthy over time** — curation prevents the stale-wiki death spiral.

“ *Honest caveat:* it only works **with** the discipline. The rules are the product, not the files. ”

Why it generalizes

- Nothing here is tool-specific. It's **caching, database indexing, lazy-loading, and good onboarding docs** — applied to an agent's working memory.
- Transferable to **any agentic AI workflow, any team**.

"Here's the pattern — happy to go deeper on the mechanics."